Medical Symptoms Checker

A Project Review Paper

Maulana Abul Kalam Azad University of Technology
June 2026

By

Moumita Paul
Registration No: 231000110287 of 2023-24
Examination Roll No: 10000123043

Under the Guidance of

Dr. Pradyut Sarkar
Associate Professor
Department of Computer Science & Engineering,
Maulana Abul Kalam Azad University of Technology,
Kolkata-741249, W.B., India

ACKNOWLEDGEMENTS

First and foremost, I would like to express my earnest gratitude and heartfelt indebtedness to my advisor, Dr.Pradyut Sarkar, Department of Computer Science & Engineering, for the privilege and the pleasure, of allowing me to work under him toward my Degree of Bachelor of Computer Science & Engineering. This work would not have materialized but for his whole-hearted help and support. Working under him has been a great experience. I sincerely thank my supervisor, particularly for all the faith he had in me. I am thankful to Dr.Pradyut Sarkar. I would also like to show my gratitude to the respected professors of the Department of Computer Science & Engineering for their constant guidance and valuable advice.

Date:
Place:

Moumita Paul
Examination Roll No: 10000123043
Maulana Abul Kalam Azad University of Technology

Table of Contents

Chapter 1	Page No
• 1.1 Abstract	4
• 1.2 Keywords	5
• 1.3 Introduction	5-6
• 1.4 Background	6-7
• 1.5 Scope of the review work	7
Chapter 2	
• 2.1 Literature Survey	8
• 2.2 Paper Comparison	9-11
○ 2.2.1 Summary of the research paper comparisons	11-12
○ 2.2.2 Challenges and Problems Identified in Existing Research	12-13
○ 2.2.3 Research Gaps in Symptom-Based Disease Prediction	13-14
○ 2.2.4 Limitations of this Review	14
Chapter 3	
• 3.1 Literature Search Strategy and Selection Criteria	14-15
• 3.2 Dataset Characteristics and Preprocessing	15
• 3.3 Model Selection and Optimization Techniques	15
• 3.4 Evaluation Metrics and Performance Analysis	15
Chapter 4	
• 4.1 Future Scope and Conclusion	16-18
• 4.2 References	18

Chapter 1

1.1 Abstract

The integration of Artificial Intelligence (AI) and Machine Learning (ML) into modern healthcare is reshaping how we approach early diagnosis, particularly through the use of medical symptom checkers. By evaluating patient-reported symptoms, these systems predict potential underlying diseases and actively support clinical decision-making. The need for this technology is particularly pressing because traditional diagnostic pathways rely heavily on expert medical knowledge, extensive physical evaluations, and immediate access to clinical facilities. For patients in underdeveloped or resource-constrained regions, these requirements often result in severe diagnostic delays. Researchers are increasingly exploring machine learning algorithms to develop automated, highly accessible disease prediction systems that can operate outside the traditional hospital setting. This review provides a comprehensive analysis of recent studies on symptom-based disease prediction and medical symptom checker systems published between 2024 and 2025. The surveyed research evaluates a wide variety of machine learning approaches, ranging from traditional models like Random Forest, K-Nearest Neighbors (KNN), and Support Vector Machines (SVM), to more complex architectures such as Gradient Boosting, Neural Networks, and ensemble learning methods. The review covers the emerging use of Large Language Models (LLMs) in clinical settings. These distinct approaches are systematically compared based on their underlying algorithms, the specific datasets utilized, their prediction accuracy, and the unique healthcare challenges they aim to solve. Across the selected literature, reported accuracies range from approximately 84% to 99%, demonstrating the rapidly growing effectiveness of AI-driven diagnostic systems. However, the review further examines several major challenges consistently identified in existing research that limit real-world deployment. Key roadblocks include severe class imbalance in medical datasets, limited dataset diversity, and the difficulty models face when processing overlapping disease symptoms. Beyond the data itself, there are significant hurdles related to computational complexity, concerns regarding patient safety and privacy, and a lack of explainability when models function without transparent reasoning. To highlight potential solutions to these limitations, the paper explores recent developments in Explainable Artificial Intelligence (XAI), medicine recommendation systems, and conversational AI-based diagnostic assistants. The findings indicate that while current symptom checkers have achieved remarkable predictive performance, they require further improvements in algorithmic transparency, multimodal data integration, and real-world clinical validation. This review provides a consolidated understanding of the current state of symptom-based disease prediction and highlights promising directions for future research in intelligent healthcare systems.

1.2 Keywords

Symptom-Based Disease Prediction, Medical Symptom Checker, Machine Learning, Artificial Intelligence, Explainable AI (XAI), Clinical Decision Support Systems, Healthcare Informatics, Large Language Models.

1.3 Introduction

Today's global healthcare infrastructure is under immense strain. As populations continue to grow, medical systems are increasingly burdened by rising treatment costs and a critical shortage of healthcare professionals. These systemic pressures inevitably lead to unequal access to medical services. According to global healthcare reports, a substantial portion of the world's population remains cut off from essential clinical facilities, a crisis that is especially pronounced in rural and underdeveloped regions. In this landscape, early disease diagnosis is arguably the single most important factor in improving patient outcomes, lowering mortality rates, and curbing overall medical expenses. However, traditional diagnostic pathways are not always feasible. They rely heavily on specialized medical expertise, extensive laboratory testing, and significant time commitments—resources that are simply out of reach for many individuals requiring timely care. To bridge this widening diagnostic gap, researchers have turned to the capabilities of Artificial Intelligence (AI) and Machine Learning (ML). These technologies have opened up entirely new avenues for designing intelligent healthcare frameworks that can support both patients and overworked clinical staff. Within this space, medical symptom checkers have emerged as a particularly promising application. The premise is straightforward but powerful: users input their symptoms, and the system cross-references these inputs against complex patterns learned from extensive medical datasets to predict potential illnesses. By functioning as accessible, preliminary screening tools, these symptom checkers help facilitate early clinical intervention, empower patients to make informed healthcare decisions, and alleviate some of the unnecessary diagnostic burden currently placed on hospitals and clinics.The technical evolution of these platforms has been rapid. Over the last few years, the research community has proposed a wide array of symptom-based prediction models powered by diverse machine learning architectures. These include foundational algorithms like Random Forest, Decision Trees, K-Nearest Neighbors (KNN), Naïve Bayes, Support Vector Machines (SVM), Gradient Boosting, and Artificial Neural Networks. The results have been highly encouraging, with multiple studies reporting prediction accuracies well over 90%, firmly establishing the viability of AI-assisted medical triage. More recently, the scope of this research has broadened significantly. The field is now pushing the boundaries of what these systems can achieve by integrating advanced ensemble learning techniques, Natural Language Processing (NLP) for better symptom parsing, integrated medicine recommendation modules, Explainable AI (XAI) frameworks, and Large Language Models (LLMs) capable of executing sophisticated clinical reasoning tasks. Yet, despite this impressive technological momentum, the transition from academic research to practical, real-world deployment is hindered by several persistent challenges. A primary issue is that many models are trained on narrow or limited datasets, which severely restricts their generalizability.

Algorithms frequently struggle to confidently differentiate between illnesses that present with highly similar, overlapping symptoms. Compounding this is the "black-box" nature of many high-performing models; they can provide a highly accurate prediction but often offer little to no explanation as to the underlying reasoning, which creates a major barrier to clinical trust. Other significant obstacles include inherent data imbalances, the ambiguity of poorly described, self-reported user symptoms, stringent patient privacy concerns, and a noticeable lack of rigorous, real-world clinical validation trials. Given this context, this review paper aims to critically examine and compare the most recent developments in symptom-based disease prediction and automated medical symptom checkers. By systematically analyzing contemporary literature, this study identifies the most frequently deployed machine learning algorithms, benchmarks their reported performance, and casts a spotlight on the critical limitations currently holding the technology back. The paper explores the emerging technological trends that are expected to define the next generation of intelligent clinical support. The objective of this review is to provide a thorough, consolidated understanding of exactly how artificial intelligence is reshaping the landscape of disease prediction and clinical decision-making.

1.4 Background

The process of transforming a patient's medical symptoms into accurate medical diagnosis has been a cornerstone for the advancement of healthcare practices for over a century. In 1923, Dr. Richard Cabot introduced the Clinicopathological Conferences (CPCs) and since then expert physicians have been using them for making curated case records for clinical reasoning and benchmark the evolution of diagnostic technology [1]. Although CPCs have been used for decades their use has considerable limitations because the CPCs are pedagogical in nature; despite being real cases, they are highly curated and information-dense, intended to demonstrate expert reasoning rather than real-life diagnosis [1]. As such, incautious uses of the CPCs can lead to overstating the capabilities of AI systems in real-life diagnosis [1]. This disconnects between curated academic data and everyday clinical practice becomes particularly significant given the current global healthcare challenges. The World Health Organisation (WHO) reports that about half of the world's population does not have access to basic healthcare services, with the gaps being most noticeable in areas with few resources. While traditional diagnosis depends on manual symptom analysis by experts, this method is frequently expensive, time consuming and prone to human error [9]. Reliance on unverified "online symptom checkers" or "Dr. Google" frequently leads to unreliable health advice and erroneous results, creating an urgent need for more credible, accurate diagnostic tools. Machine learning (ML) and artificial intelligence (AI) address these limitations by analyzing large volumes of data to identify patterns and predict illnesses based on reported symptoms. Various machine learning techniques, including decision-tree-based models, support vector machines, nearest-neighbor methods, and deep learning approaches, have been explored to improve disease prediction accuracy. Advanced approaches now integrate Convolutional Neural Networks (CNNs) and word embedding to analyze symptom similarity, or extract semantic information from unstructured clinical notes—which often contain more patient-specific insights than structured Electronic Health Records (EHRs)—to improve predictive accuracy. The primary objective of these symptom-based prediction systems is to facilitate early intervention, which significantly enhances patient outcomes and reduces the risks of misdiagnosis or treatment delays. By providing both healthcare professionals and individuals with reliable, data-driven insights, these systems transform healthcare from a purely curative model into a proactive, preventive, and personalized service. At the end, these technologies aim to alleviate the burden on the pharmaceutical and clinical landscapes by turning massive amounts of raw patient data into actionable, meaningful diagnostic guidance.

1.5 Scope of the review work

This review focuses on analyzing the latest research surrounding medical symptom checkers and AI-driven disease prediction. Specifically, the scope centers on systems leveraging Artificial Intelligence, Machine Learning, and Deep Learning to predict illnesses based entirely on patient-reported symptoms. To ensure the findings are current and relevant, the analysis primarily investigates literature published between 2024 and 2025, capturing a wide spectrum of modern diagnostic approaches. Within this timeframe, the paper systematically examines a diverse set of machine learning architectures. This includes foundational algorithms like Random Forest, K-Nearest Neighbors (KNN), and Support Vector Machines (SVM), alongside more complex frameworks such as Gradient Boosting, Artificial Neural Networks, ensemble learning models, and emerging Large Language Models (LLMs). The selected studies are critically compared across several dimensions: their underlying methodologies, the datasets they rely on, their reported prediction accuracy, and their specific advantages and clinical limitations. The review pays special attention to specialized healthcare applications, particularly systems built for multi-disease prediction, initial disease screening, clinical decision support, automated medicine recommendation, and user-facing web platforms. Beyond evaluating model performance, this work also thoroughly explores the practical roadblocks currently limiting these systems. It addresses critical research challenges such as severe class imbalances in medical data, limitations in feature selection, the difficulty of distinguishing between diseases with overlapping symptoms, and the high computational complexity of advanced models. It also highlights urgent, real-world concerns regarding patient safety, data privacy, and the inherent lack of model interpretability. To map out potential solutions to these hurdles, the paper discusses the integration of emerging technologies like Explainable Artificial Intelligence (XAI), Natural Language Processing (NLP), conversational AI, and intelligent diagnostic assistants. To maintain a sharp focus, the boundaries of this review are strictly limited to symptom-based prediction models. It explicitly excludes detailed analyses of purely image-based diagnostics, genomic prediction models, or frameworks that rely solely on laboratory data. The primary objective is to deliver a comprehensive overview of the field's current advancements, pinpoint existing research gaps, and highlight the future directions necessary to build medical symptom checkers that are accurate, reliable, highly explainable, and ready for real-world clinical application.

Chapter 2

2.1 Literature Survey

Early advancements in disease prediction relied heavily on traditional supervised learning algorithms. Studies have consistently demonstrated the effectiveness of Decision Trees, Random Forest, Naïve Bayes, and K-Nearest Neighbors (KNN) in processing patient symptoms [3, 5, 9]. Research utilizing datasets of up to 4,920 patient records has shown that these ensemble methods can achieve high classification accuracy, often ranging between 95% and 98% [3, 9]. Specifically, comparative analyses have highlighted that while algorithms like Naïve Bayes are effective for symptom-based prediction, SVM often excels in specialized detection for conditions such as kidney disease and Parkinson’s [5, 6].

To overcome the limitations of basic models—such as low diagnostic relevance and training complexity—recent research has shifted toward neural models and hybrid architectures. Nesterov et al. [3] introduced a neural model with logic regularization that outperforms traditional reinforcement learning methods, particularly in large, sparse symptom spaces. Similarly, the use of deep learning has expanded beyond static symptoms to medical imaging and time-correlated sensor data.

Beyond mere diagnosis, modern literature explores the application of AI in patient management and medication recommendations. Research has moved toward "case-similarity" retrieval systems that assist physicians in sharing knowledge and optimizing care for chronic conditions like Type 2 diabetes [6]. Semantic-enabled services like GalenOWL have been developed to provide real-time interaction discovery by translating medical data into ontologies [6]. These tools are vital in mitigating medication errors, which have been linked to limited clinical experience in up to 42% of cases [6].

Despite high accuracy rates, significant limitations persist. A primary concern is the "black-box" nature of many high-performing models, which hampers clinical trust. There is an increasing push for Explainable AI (XAI) frameworks, such as SHAP and LIME, to provide transparency in model decisions [8, 11]. Many existing models struggle with data imbalance and the clinical reality of overlapping symptoms across different diseases [7, 9, 11]. Recent studies emphasize that for a model to be viable in real-world applications, it must move away from single-disease focus and demonstrate generalizability across a wide array of medical conditions and patient demographics [8, 9, 11].

However, a comprehensive symptom checker capable of accurately predicting multiple diseases across diverse populations while maintaining transparency and clinical interpretability remains an open research challenge.

2.2 Paper Comparison

Here is the comparison of the 12 research papers dated from 2024-2025 I used in this review work:

Year	Title	Accuracy	Algorithm Used	Problem Addressed
2025	Advancing Medical Artificial Intelligence Using a Century of Cases	84%	Large language models (OpenAI GPT-4o "o3", Google Gemini 2.5 Pro) plus an "AI discussant" agent ("Dr. CaBot")	Differential diagnosis on historical NEJM clinical cases (CPC-Bench tasks)
2025	SYMPTOM-BASED DISEASE PREDICTION USING MACHINE LEARNING	99%	Random Forest Classifier and Multi-Layer Perceptron	Predict diseases from symptom data to improve diagnostic access (e.g. in low-resource settings)
2024	Optimized Machine Learning Classifiers for Symptom-Based Disease Screening	98%	K-Nearest Neighbors (k=2) (best performer) alongside SVM, Random Forest, and Neural Network	Symptom-based disease screening to alleviate healthcare system overload (serve as an initial triage)
2024	Enhancing Disease Diagnosis with Machine Learning Based Symptom Prediction	91-95%	Gradient Boosting, Random Forest (with RFE or CHI2 selection)	Traditional feature pruning pipelines completely discard subtle, low-frequency symptoms critical for rare conditions.
2024	Human Disease Prediction Using Machine Learning	87-93%	Convolutional Neural Networks (CNN), Random Forest, KNN	Rigid reliance on structured input values; black-box decision models limit adoption by working physicians.
2025	Medicine Recommendation	96% (Random	Random Forest, Support Vector Classifier (SVC), and	Recommend medicines for

	System for Diseases Using Machine Learning	Forest) and 92% (SVC).	NLP models for symptom parsing.	diagnosed diseases (ML-driven decision support)
2025	Optimizing Symptom Prediction in Healthcare: An Integrative Machine Learning Approach to Regression and Classification Models	91.84%	Integrated Ensemble Models (primarily Random Forest Classifier), Recursive Feature Elimination, and hybrid resampling tactics (SMOTE/undersampling).	Overcoming extreme class imbalances inherent to rare disease conditions and rendering black-box ensemble decisions clear enough for clinical compliance.
2024	Predicting Multiple Diseases Using Machine Learning: A Symptom-Based Model for Medical Diagnosis	90.74%	K-Nearest Neighbors (KNN), Decision Trees, and specialized Bagging classifiers.	Predict multiple diseases from a patient's symptoms
2024	Prediction System for Symptoms Driven Disease Diagnosis	95%	Random Forest Classifier and Naïve Bayes.	Users entering vague, qualitative, or mistranslated symptoms which cause system failures or erroneous outputs.
2024	Symptom-Based Disease Prediction Using Machine Learning: A Web Application Approach	92-95%	SVM, Naive Bayes, Random Forest	Web-based tool for symptom-driven disease prediction (real-time user-facing application)
2024	Symptoms-based Classification of Disease Using Various ML Algorithms and Interpretation using LIME and SHAP KERNELS	95-96%	Gradient Boosting, Random Forest, LIME (Local Interpretable Model-agnostic Explanations), and SHAP (SHapley Additive exPlanations) kernels.	High computational footprint and latency during the real-time calculation of SHAP values for complex patient profiles.

2024	Disease Prediction System	94-95%	Supervised tree algorithms combined with basic Natural Language Processing (NLP) interfaces.	The potential for user self-misdiagnosis due to a lack of physical vitals validation (e.g., blood pressure, heart rate) within software systems.

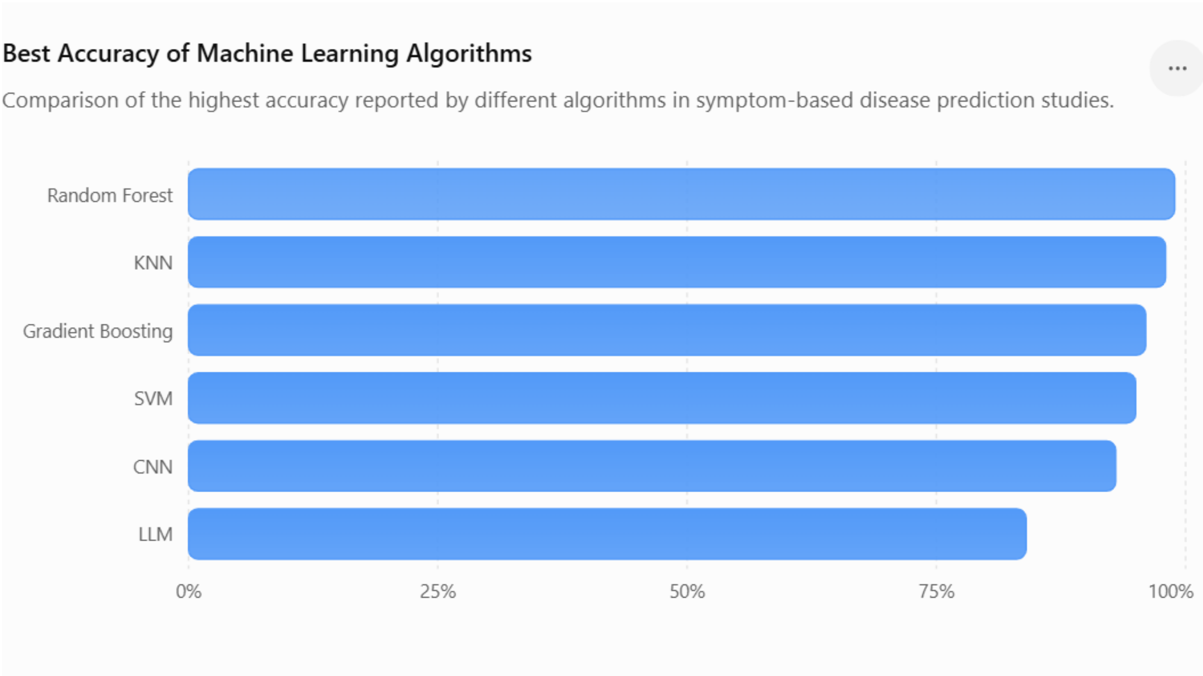

The values represent the highest reported accuracy for each algorithm among the reviewed studies.

2.2.1 Summary of the research paper comparisons

The reviewed studies show that machine learning and artificial intelligence have significantly improved the effectiveness of medical symptom checkers and disease prediction systems. Most of the research focuses on predicting diseases from patient symptoms, improving diagnostic accuracy, and providing early healthcare assistance to users. Studies [2], [3], [9], and [10] demonstrate that traditional machine learning algorithms such as Random Forest, K-Nearest Neighbors (KNN), Support Vector Machine (SVM), and Naïve Bayes can achieve high prediction accuracy when trained on symptom datasets. Among these, Random Forest emerged as one of the most frequently used algorithms due to its robustness and consistent performance. The highest reported accuracy of 99% was achieved in [2], while [3] reported an accuracy of 98% using an optimized KNN model. Several researchers have attempted to address limitations of existing disease prediction systems. Study [4] focused on preserving important low-frequency symptoms that are often ignored during feature selection, thereby improving the diagnosis of rare diseases. Similarly, study [7] addressed the challenge of class imbalance in healthcare datasets by integrating ensemble learning techniques with data resampling methods. These approaches contributed to improved model performance and reliability. Interpretability has also become an important area of research. Study [11] incorporated explainable AI techniques such as LIME and SHAP to provide clear explanations for disease predictions. This helps healthcare professionals understand the reasoning behind model decisions and increases trust in AI-based systems. However, the study also highlighted the computational overhead associated with generating explanations for complex cases. Some studies expanded the scope of symptom checkers beyond disease prediction. Study [6] introduced a medicine recommendation system that combines disease prediction with treatment suggestions, while [10] developed a web-based platform that enables users to obtain real-time predictions through an accessible interface. These systems demonstrate the growing trend toward user-centered healthcare applications. Unlike traditional machine learning approaches, study [1] explored the use of large language models for medical diagnosis. Although its reported accuracy was lower than many specialized machine learning models, the research highlighted the potential of advanced AI systems in clinical reasoning and diagnostic support. This indicates a shift toward integrating generative AI technologies into healthcare applications. In all things considered, the reviewed literature indicates that Random Forest, KNN, Gradient Boosting, and ensemble learning techniques consistently achieve high prediction accuracy, generally ranging from 90% to 99%. Recent research is increasingly focused not only on improving predictive performance but also on enhancing model interpretability, handling imbalanced datasets, supporting medicine recommendations, and developing user-friendly healthcare systems. These advancements suggest that future medical symptom checkers will become more accurate, explainable, and accessible, thereby contributing to improved healthcare decision-making and early disease detection.

2.2.2 Challenges and Problems Identified in Existing Research

The reviewed studies reveal several challenges that continue to affect the effectiveness and reliability of medical symptom checker systems. Although significant improvements have been achieved through machine learning and artificial intelligence, many limitations remain unresolved. One of the most common challenges is the handling of class imbalance in medical datasets. Rare diseases often have fewer training samples compared to common diseases, causing prediction models to favor frequently occurring conditions and reducing their ability to identify uncommon illnesses accurately [7]. Another major issue is the loss of important symptom information during feature selection. Many machine learning models remove symptoms that appear infrequently in the dataset, even though these symptoms may be crucial for diagnosing rare or complex diseases. This can lead to incorrect predictions and reduced diagnostic reliability [4]. Several studies also highlight concerns regarding model interpretability and transparency. While advanced algorithms such as Random Forest, Neural Networks, and ensemble models achieve high accuracy, healthcare professionals often find it difficult to understand how these systems arrive at their predictions. The lack of explainability can limit clinical acceptance and trust in AI-based diagnosis [5], [11]. The quality of predictions is also affected by inaccurate symptom input from users. Patients may describe symptoms vaguely, use non-medical terminology, or provide incomplete information. Such inconsistencies can reduce system accuracy and lead to incorrect disease predictions [9]. Another challenge is the absence of clinical and physiological data in many symptom checker systems. Most reviewed models rely solely on symptoms and do not consider vital signs, laboratory reports, medical history, or imaging results. As a result, the predictions may not fully reflect the patient's actual health condition [12]. Computational complexity is another concern, particularly in systems using explainable AI techniques such as LIME and SHAP. Although these methods improve transparency, they increase processing time and resource consumption, making real-time deployment more challenging [11]. Web-based and self-diagnosis systems may encourage self-medication and misinterpretation of results if users rely solely on automated predictions without consulting healthcare professionals. This raises concerns about patient safety and responsible use of medical AI tools [6], [10]. To sum it up, the literature indicates that future medical symptom checker systems must focus on improving dataset quality, handling class imbalance, enhancing model explainability, incorporating additional clinical information, and ensuring safe and responsible use. Addressing these challenges will help develop more reliable, accurate, and trustworthy healthcare support systems.

2.2.3 Research Gaps in Symptom-Based Disease Prediction

While the reviewed literature collectively demonstrates that machine learning and AI-driven symptom checkers have achieved impressive predictive accuracy—often reaching or exceeding 95%—a critical analysis reveals several persistent gaps that limit their transition from academic research to real-world clinical deployment. First, there is a significant disconnect between high-performing models and clinical transparency. Many of the studies reviewed rely on "black-box" architectures, such as complex ensemble models or unoptimized neural networks, which provide highly accurate predictions but offer no explanation for their reasoning. Despite the recent introduction of interpretability tools like LIME and SHAP in studies like [11], these remain secondary additions rather than fundamental components of the model design, often resulting in high computational latency that makes them difficult to use in real-time settings. Second, most existing systems are constrained by a narrow diagnostic focus. A large majority of models are trained on curated, limited datasets that do not account for the messy reality of patient-reported symptoms. We see a recurring failure in models to differentiate between diseases that share overlapping symptoms—a common challenge in clinical practice that current systems, which often rely on simplified, structured binary inputs, struggle to resolve. These models frequently overlook physiological data, such as vital signs or patient history, which are essential for an accurate diagnosis but are largely absent from the current symptom-checker framework. Third, there is a notable lack of real-world validation. Almost all the reviewed research relies on controlled, offline experiments rather than prospective clinical trials. While studies like [1] have begun to explore the clinical reasoning capabilities of Large Language Models, the leap from predicting a disease in a lab-tested dataset to safely triaging a patient in a rural or resource-constrained environment remains largely unaddressed. Finally, the issue of data imbalance continues to be a bottleneck. Rare diseases are consistently underrepresented in the datasets used in these studies, leading to models that inherently prioritize common conditions while failing to provide reliable screening for rarer or more complex illnesses. Bridging these gaps—moving toward transparent, multimodal, and clinically validated frameworks—is the necessary next step for the field to evolve into a truly trustworthy diagnostic tool.

2.2.4 Limitations of this Review

While this analysis provides a comprehensive overview of the current landscape of AI-driven symptom checkers, it is important to acknowledge the inherent boundaries of this literature survey. First, the scope of this review is strictly confined to research published between 2024 and 2025. Given the rapid velocity of advancements in Large Language Models (LLMs) and deep learning architectures, this analysis represents a snapshot in time. It may not encompass emergent breakthroughs or methodological shifts that have occurred outside of this specific publication window. Second, the review methodology explicitly prioritizes symptom-based prediction frameworks. By design, this excludes other vital branches of medical AI, such as purely image-based diagnostics (e.g., radiology or CT scan analysis) and genomic prediction models. While this focus ensures consistency in the analyzed data, it means the findings should not be interpreted as a universal assessment of all AI applications in healthcare, but rather as a specific exploration of triage and preliminary screening tools. Finally, because this is a secondary analysis based on existing academic studies, the review's conclusions are naturally tethered to the constraints of the source material. Many of the primary papers rely on curated, structured datasets or offline experimental setups rather than longitudinal clinical trials. As such, this review reflects the current state of academic experimentation. Readers should view the reported high-accuracy metrics with the understanding that they stem from controlled, lab-based environments and have yet to be universally validated across diverse, heterogeneous, and unpredictable real-world patient populations.

Chapter 3

3 Methodology

3.1 Literature Search Strategy and Selection Criteria

To synthesize the latest advancements in symptom-based disease prediction, a systematic review framework was established. The primary search targeted literature published between 2024 and 2025 across major academic databases. The core inclusion criteria required studies to demonstrate empirical results on multi-disease prediction utilizing supervised or ensemble machine learning techniques. A total of twelve studies were selected for this review. However, five highly relevant studies ([3], [4], [7], [8], and [9]) were analyzed in greater methodological detail to examine dataset characteristics, preprocessing strategies, model optimization techniques, and evaluation metrics. Theoretical proposals lacking experimental validation and models relying strictly on non-symptom medical imaging were excluded from this analysis.

3.2 Dataset Characteristics and Preprocessing

The majority of the selected studies relied on structured symptom-disease matrices to train their models. For instance, multiple researchers evaluated frameworks using a standardized dataset encompassing 4,920 patient records spanning 132 binary symptom features to classify 41 distinct diseases. A recurrent theme across the literature was the critical need for robust data preprocessing. The authors in [8] explicitly addressed inherent dataset disparities by implementing resampling tactics to manage the class imbalances that occur when specific rare diseases have significantly fewer data samples. Processing unstructured and often vague patient input was another key focus. The study in [7] developed a deterministic symptom indexing approach specifically designed to convert free-text symptom inputs into fixed-length binary feature vectors. To ensure that critical diagnostic data was not lost during feature pruning, the methodology in [4] utilized sentence similarity models, syntactic trees, and word embeddings to extract vital keyword information from natural patient descriptions.

3.3 Model Selection and Optimization Techniques

Across the five studies, researchers evaluated a broad spectrum of classification algorithms. The research in [3] reported a rigorous three-phase optimization pipeline consisting of initial data preprocessing, hyperparameter grid search variations, and final model filtering. This intensive pipeline identified K-Nearest Neighbors (KNN with 2 neighbors) as the optimal model for rapid initial patient screening. In contrast, for environments requiring multi-label classification for concurrent diseases, the framework in [8] optimized ensemble methods and addressed overlapping health conditions using a One vs. All algorithmic approach. The architecture proposed in [4] pushed the complexity further by integrating Convolutional Neural Networks (CNNs) directly into the computation process to improve early disease intervention based on nuanced symptom similarities. The methodological analysis of these architectures considered their capacity for secure deployment; specifically, assessing how such models might integrate with decentralized federated learning frameworks to ensure stringent patient data privacy during the predictive diagnostic process.

3.4 Evaluation Metrics and Performance Analysis

To benchmark these varied models, standard metrics including accuracy, precision, recall, and F1-score were heavily utilized. The optimized KNN model in [3] achieved an F1-score and accuracy exceeding 98%. Meanwhile, multi-class classifiers, such as Support Vector Classifiers (SVC) with linear kernels, were trained on curated symptom mappings to prioritize fast, reliable inference in other works. When dealing with disproportionate datasets, the evaluation in [8] prioritized the F1-Score over raw accuracy, as its harmonic mean of precision and recall is significantly more beneficial for evaluating imbalanced medical classes. The comparative analysis focused not only on aggregate accuracy but also on execution time and computational footprint, recognizing that a viable clinical screening tool must operate efficiently in real-world settings.

Chapter 4

4.1 Future Scope and Conclusion

The field of medical symptom checkers and symptom-based disease prediction systems is rapidly evolving due to advancements in Artificial Intelligence (AI), Machine Learning (ML), Deep Learning (DL), and Large Language Models (LLMs). Although recent studies have reported impressive prediction accuracies exceeding 90%, several challenges remain regarding real-world deployment, interpretability, scalability, and clinical reliability. One major future direction is the integration of Large Language Models (LLMs) into symptom-checking platforms. The study *Advancing Medical Artificial Intelligence Using a Century of Cases (2025)* demonstrated that advanced AI systems can perform complex differential diagnosis and clinical reasoning tasks with performance comparable to or even exceeding human physicians in specific scenarios. Future symptom checkers may employ LLMs to conduct natural conversations with patients, ask follow-up questions dynamically, interpret symptom descriptions in everyday language, and generate clinically meaningful explanations for predicted diseases. Such conversational AI systems can significantly improve accessibility and user experience. Another important area of development is the creation of multimodal diagnostic systems. Most current symptom checkers rely solely on textual symptom input. Future systems can combine symptom data with medical images, laboratory reports, wearable sensor data, electronic health records (EHRs), and physiological measurements such as heart rate, blood pressure, oxygen saturation, and body temperature. By integrating multiple sources of patient information, these systems can achieve higher diagnostic accuracy and reduce the risk of misclassification caused by overlapping symptoms.The adoption of Explainable Artificial Intelligence (XAI) is expected to become a fundamental requirement for clinical acceptance. Several studies reviewed in this paper highlighted the black-box nature of high-performing machine learning models as a major limitation. Future systems should incorporate advanced explainability techniques such as SHAP, LIME, attention visualization, and rule-based explanations to provide transparent reasoning behind predictions. Such transparency will help healthcare professionals understand, validate, and trust AI-assisted diagnoses. Future research should also focus on privacy-preserving machine learning frameworks. Medical data are highly sensitive, and concerns regarding patient confidentiality continue to limit data sharing among healthcare institutions. Techniques such as Federated Learning, Differential Privacy, Secure Multi-Party Computation, and Blockchain-based healthcare architectures can enable collaborative model training while protecting patient information. This would allow models to learn from larger and more diverse datasets without compromising privacy.

A significant limitation of current systems is their dependence on relatively small and structured datasets. Future studies should develop large-scale, diverse, and multilingual symptom datasets that better represent real-world patient populations. Most existing models are trained on datasets containing approximately 41 diseases and 130–140 symptoms. Future systems should expand their coverage to include rare diseases, emerging diseases, chronic conditions, pediatric disorders, geriatric diseases, and region-specific health conditions. This will improve model generalizability and clinical applicability across different demographic groups. Another promising direction is the incorporation of personalized medicine into symptom-checking systems. Future models can consider patient-specific factors such as age, gender, family history, lifestyle habits, genetic predispositions, environmental exposures, and previous medical records. Personalized prediction models can provide more accurate and individualized disease assessments than generalized population-based approaches.Future symptom checkers are also expected to play a critical role in telemedicine and remote healthcare services. With increasing demand for virtual healthcare, AI-powered symptom checkers can serve as intelligent triage systems that prioritize urgent cases, recommend appropriate medical specialists, and guide patients toward suitable healthcare resources. This capability can reduce hospital overcrowding and improve healthcare accessibility, particularly in rural and underserved regions. Recent studies have further demonstrated the effectiveness of ensemble learning and hybrid architectures. Therefore, future research should explore hybrid AI models that combine Random Forests, Deep Neural Networks, Transformer-based architectures, Knowledge Graphs, and Expert Systems. Such integrated approaches may leverage the strengths of multiple techniques while minimizing their individual limitations.

At the last but not the least, future systems should undergo extensive clinical validation and real-world testing. Most current studies evaluate performance using controlled datasets and offline experiments. Large-scale prospective clinical trials involving hospitals, healthcare professionals, and diverse patient populations are necessary to assess reliability, safety, fairness, and effectiveness under real-world conditions. Regulatory compliance and ethical AI guidelines will also become increasingly important as these systems move toward clinical deployment.

This review paper examined recent developments in medical symptom checkers and symptom-based disease prediction systems utilizing Machine Learning, Deep Learning, and Artificial Intelligence techniques. The analysis of selected studies published between 2024 and 2025 reveals that symptom-based disease prediction has emerged as a promising approach for supporting early diagnosis, improving healthcare accessibility, and reducing diagnostic delays. The reviewed literature demonstrates that traditional machine learning algorithms such as Random Forest, K-Nearest Neighbors (KNN), Support Vector Machines (SVM), Naïve Bayes, and Decision Trees continue to achieve strong predictive performance. Among these approaches, Random Forest-based models consistently demonstrated superior accuracy, robustness, and scalability across multiple disease prediction tasks. Several studies reported accuracy levels exceeding 95%, with some models achieving up to 99% prediction accuracy under controlled experimental conditions. Recent advancements in deep learning, ensemble learning, and feature optimization techniques have further enhanced diagnostic performance. The literature also indicates a gradual transition from conventional machine learning systems toward more sophisticated AI-driven healthcare solutions. Explainable AI techniques, Natural Language Processing, intelligent recommendation systems, and Large Language Models are increasingly being incorporated into symptom-checking platforms to improve interpretability, user interaction, and clinical decision support. The emergence of systems such as Dr. CaBot and CPC-Bench highlights the growing capability of AI to perform complex clinical reasoning and differential diagnosis tasks. Despite these encouraging developments, several challenges remain unresolved. Many existing models rely on limited datasets, struggle with overlapping symptoms among multiple diseases, exhibit reduced generalizability across diverse populations, and often function as black-box systems that provide limited insight into their decision-making processes. Issues related to patient privacy, ethical AI deployment, regulatory approval, and real-world clinical validation continue to hinder widespread adoption.

In conclusion, medical symptom checkers have significant potential to transform healthcare by enabling faster, more accessible, and data-driven disease prediction. While current systems demonstrate impressive accuracy, future advancements should prioritize explainability, privacy preservation, multimodal data integration, personalized healthcare, and real-world clinical validation. By addressing these challenges, next-generation symptom-checking systems can evolve from simple prediction tools into reliable clinical decision-support systems that assist both patients and healthcare professionals in delivering efficient, accurate, and personalized healthcare services.

4.2 References:

[1] T. A. Buckley, R. Conci, P. G. Schlegel, et al., “Advancing Medical Artificial Intelligence Using a Century of Cases,” 2025.

[2] V. Pal et al., “Symptom-Based Disease Prediction Using Machine Learning: A Web Application Approach,” 2024.

[3] A. Fuster-Palà et al., “Optimized Machine Learning Classifiers for Symptom-Based Disease Screening,” 2024.

[4] P. Subramaniam et al., “Enhancing Disease Diagnosis with Machine Learning Based Symptom Prediction,” 2024.

[5] K. Gaurav et al., “Human Disease Prediction Using Machine Learning Techniques and Real-Life Parameters,” 2024.

[6] H. S. S. Kumar et al., “Medicine Recommendation System for Diseases Using Machine Learning,” 2025.

[7] N. R. Goddilla and A. Sivasangari, “Optimizing Symptom Prediction in Healthcare: An Integrative Machine Learning Approach to Regression and Classification Models,” 2025.

[8] B. Makhkamov et al., “Predicting Multiple Diseases Using Machine Learning: A Symptom-Based Model for Medical Diagnosis,” 2024.

[9] A. Sharma et al., “Prediction System for Symptoms Driven Disease Diagnosis,” 2024.

[10] R. Sood and V. Sharma, “Symptom-Based Disease Prediction Using Machine Learning,” 2024.

[11] K. J. Soujanya et al., “Symptoms-Based Classification of Disease Using Various ML Algorithms and Interpretation Using LIME and SHAP Kernels,” 2024.

[12] R. Khiratkar et al., “Disease Prediction System,” 2024.